

BUILT ON BENTO × ANAKIN

HATE ON YOUR FRIENDS. ON PURPOSE.

A social accountability app where your friends put real testnet credits on whether you'll actually do the thing. An AI oracle settles it. No self-reporting.

HATE (NO)

BELIEVE (YES)

THE PROBLEM

HABIT APPS ASK YOU TO GRADE YOUR OWN HOMEWORK.

SELF-REPORTED

You tap "done." Nobody checks. The streak is a number you award yourself.

NO STAKES

Breaking a streak costs nothing, so the streak stops meaning anything by day four.

NO AUDIENCE

Nobody in your group chat knows or cares that you skipped the gym again.

SpiteBet fixes all three at once: your friends have money on it, they can see exactly who bet against you, and a machine checks the receipts.

THE LOOP

FOUR STEPS, AND THEN IT RUNS ITSELF.

1

POST A TRAP

"I'll make 5 commits before 5pm." Attach your GitHub handle, pick a deadline. This mints a real prediction market on Bento.

2

FRIENDS SWIPE

Left to hate, right to believe. Each bet is signed by that person's own wallet, so the vote count is real – and public.

3

THE ORACLE JUDGES

At the deadline it pulls your commits and an AI decides whether they're substantive. Ten README typo commits do not pass.

4

POOL SETTLES

Winning side takes the credits. Everyone who bet opens the market and sees the verdict, the evidence, and their payout.

NOT A MOCKUP

EVERY NUMBER IN THE APP COMES FROM BENTO.

6

MARKETS MINTED ON-CHAIN

4

REAL TESTNET ACCOUNTS

3

DISTINCT VOTERS ON ONE MARKET

```
POST /bento/user/duels/create      → 201
  duelId  9f42e01312e06bb657e4cd59779c40fe ...
  txHash  0x969e27e6df9d280b065b0900dabda2a1 ...

GET  participants/9f42e013 ...
     3 voters | merest.sa.slayer, karthik, term101112
```

Balance, positions, PnL, pool split and voter headcount are all read back from Bento – nothing in the app is a local counter. A stranger (karthik) found one of our markets on the shared testnet and bet on it.

DESIGN DECISION

A VOTE IS A PERSON, NOT A POOL BALANCE.

WHAT WE SHOW

"2 haters vs 1 believer," with names attached. Pulled from Bento's participants endpoint and de-duplicated per account.

MEREST_SLAYER

TERMINATOR_43

WICKED-IRIS

WHY NOT THE POOL

Creating a market seeds *both* sides with liquidity, so raw pool numbers say "300 vs 260" before anyone has voted at all. That reads as activity that didn't happen.

Anyone holding both sides is attributed to whichever side they hold more of, so the creator's seed doesn't count as two votes.

THE ORACLE

FIVE HOPS FROM DEADLINE TO PAYOUT.

01

BENTO

Read the market, its deadline,
and the GitHub handle stored in
its tags.

02

ANAKIN WIRE

Verify the account exists via
the scraping API.

03

GITHUB

Pull every commit authored in
the exact window.

04

GROQ

Llama 3.3 judges which commits
are substantive work.

05

SETTLE

Resolve on Bento; winning side
takes the pool.

Runs entirely server-side – the Anakin and Groq keys are never
shipped to the browser. Fires automatically 15 seconds after any
deadline passes.

REAL OUTPUT

COUNTING ISN'T JUDGING.

GOAL MET

BELIEVERS TAKE THE POOL

9 commits found · 7 judged substantive

"The goal was met as 7 out of 9 commits contained substantive work."

"Looks like someone's been busy, but also likes to commit every single thought."

WHY AN AI, NOT A COUNTER

"Make 5 commits" is trivially gamed – five whitespace edits clears it. The model reads the messages and discounts padding: typo fixes, empty commits, version bumps.

The same judging layer works for any activity list, which is what makes new verifiers cheap to add.

WHERE IT STANDS

WHAT'S REAL, AND WHAT ISN'T YET.

CAPABILITY	STATUS	NOTE
Create market on Bento	LIVE	Real on-chain, confirmed txHash
Place bet / AMM pricing	LIVE	estimateBuy → placeBet
Balance & positions	LIVE	Read from portfolio API
Multi-user vote counts	LIVE	4 accounts, verified 3 voters
GitHub verification	LIVE	Commits + AI quality judging
Anakin Wire GitHub actions	BLOCKED	Upstream scraper error; we fall back
Strava / Screen Time / photo	NEXT	Selectable, not yet wired

Verification is GitHub-only today. Every other verifier is recorded and bet on, but resolves as unverifiable rather than guessing. Adding one means writing a single evidence-gathering function – the judging layer is already generic.

WHAT'S NEXT

STRAVA. SCREEN TIME. GYM CHECK-INS. PHOTOS.

The hard parts are done: markets settle on-chain, votes are attributable to real people, and an AI already knows how to tell work from padding. Every new goal type is one scraper away from working.

BUILT WITH

React 19 · TypeScript · Vite · Tailwind v4 · Framer Motion ·
Bento SDK · Anakin Wire · Groq (Llama 3.3 70B)

BECAUSE NOBODY WANTS
TO LOSE TO THEIR GROUP CHAT.

Testnet credits only. No real money, no real regrets.